

High-Performance Asynchronous Pipelines: An Overview

Steven M. Nowick

Columbia University

Montek Singh

University of North Carolina at Chapel Hill

Editor's note:

Pipelining is a key element of high-performance design. Distributed synchronization is at the same time one of the key strengths and one of the major difficulties of asynchronous pipelining. It automatically provides elasticity and on-demand power consumption. This tutorial provides an overview of the best-in-class asynchronous pipelining methods that can be used to fully exploit the advantages of this design style, covering both static and dynamic logic implementations.

—Luciano Lavagno, Politecnico di Torino

We also briefly discuss design trade-offs, performance evaluation, system-level analysis and optimization techniques, CAD tool support, testing, and recent industrial and academic applications.

Applications of pipelining in asynchronous systems

For synchronous systems, pipelining is a straightforward technique: complex

■ **ONE OF THE FOUNDATIONS** of high-performance digital system design is the use of pipelining. In synchronous systems, for several decades, pipelining has been the fundamental technique used to increase parallelism and hence boost system throughput—whether for high-performance processors, multimedia and graphics units, or signal processors.

This article provides an overview of pipelining in asynchronous, or clockless, digital systems. We do not attempt an exhaustive coverage, but rather introduce the basics of several leading representative styles. These pipelines naturally fall into two classes: those that use static logic versus those that use dynamic logic for the data path. Each class tends to use a distinct approach for its control and data storage. For static logic, we introduce the classic micropipeline of Sutherland,¹ along with two high-performance variants: Mousetrap² (which uses a standard cell design) and GasP³ (which uses a custom design). For dynamic logic, we present the classic PS0 pipeline of Williams and Horowitz,^{4,5} along with two high-performance variants: the precharge half-buffer (PCHB) pipeline⁶ (which provides greater timing robustness) and the high-capacity (HC) pipeline⁷ (which provides double the storage capacity).

function blocks are subdivided into smaller blocks, registers are inserted to separate them, and the global clock is applied to all registers. In contrast, for asynchronous systems, there is no global clock. Therefore, a protocol for the interaction of neighboring stages must be defined, as well as choices of data encoding and storage elements. In addition, an explicit distributed control structure must be designed. Together, this ensemble constitutes a template or skeleton for coordinating the blocks of a pipelined asynchronous system.

Several leading processors from the 1950s and 1960s used asynchronous circuits extensively, including the Illiac and Illiac II (University of Illinois), the Atlas and MU-5 (University of Manchester), and designs from the Macromodules project (Washington University, St. Louis).

The basic concept and design of an asynchronous pipeline were presented by David Muller in his seminal paper from 1963.⁸ Since then, asynchronous pipelines have had broad application, from the early commercial graphics and flight simulation systems of Evans & Sutherland, whose LDS-1 (Line Drawing System-1) was first shipped to Bolt, Beranek and Newman (BBN) in 1969, to the foundational

approaches of Chuck Seitz.⁹ More recently, high-performance asynchronous pipelines have been used commercially in

- Sun's UltraSparc IIIi computers for fast memory access;
- the Speedster FPGAs of Achronix Semiconductor (<http://www.achronix.com>), which, at a peak performance of 1.5 GHz, are currently claimed as the world's fastest;¹⁰ and
- the Nexus Ethernet switch chips of Fulcrum Microsystems, an asynchronous start-up company recently acquired by Intel¹¹ (<http://www.fulcrummicro.com>).

Asynchronous pipelines have also been used experimentally at IBM Research for a low-latency finite-impulse response (FIR) filter chip.¹²

Asynchronous vs. synchronous pipelines

Figure 1 shows a highly simplified view of a synchronous pipeline and an asynchronous pipeline. For the latter, **the figure doesn't show the control and storage; the focus is simply on the local interstage communication.**

Each asynchronous interstage link is a communication channel, typically including both data and control. Communication is usually bidirectional, and is implemented by a *handshaking protocol*: data (and possibly a request control signal) is sent from left to right, and an acknowledge control signal is sent from right to left.

This figure highlights the fundamental differences between synchronous and asynchronous pipelines. In synchronous pipelines, the clock advances the entire system in lock step: every data item moves to the next stage on the active clock edge. Hence, the pipeline acts as a form of synchronous shift register with computation blocks. Furthermore, each stage's critical-path delay must be less than the fixed clock period, otherwise the system will malfunction. As a result, all stages typically have roughly balanced delays.

In contrast, asynchronous pipelines have no central clock. The advance of data items is coordinated locally, on a per-stage basis. Typically, a stage N can accept a new data item if two conditions hold: its left neighbor, stage $N - 1$, is providing new data; and its right neighbor, stage $N + 1$, has stored (or finished computing on) the previous data item.




Figure 1. An abstracted view of synchronous (a) versus asynchronous (b) pipelines.

(Alternatively, the second condition is equivalent to stating that stage N has become empty.) Hence, using this local rule, data advances in a decoupled manner through the pipeline: different data items can progress at different rates. Almost all asynchronous pipelines, including those using advanced techniques, employ some form of this local rule for data movement.

The asynchronous approach has four important features in providing design flexibility and modularity. First, stages need not have equal delays. In most synchronous systems, the worst-case stage delay must be less than the clock period, and all stages operate at the same fixed rate. In contrast, in an asynchronous system, although balanced stages tend to provide the highest system performance, this balance is not a requirement for correct operation. As a result, stages of widely different static delays can be concatenated to form a working system. In addition, in some asynchronous pipeline styles, each stage may have a dynamically varying delay, such as an asynchronous adder that has data-dependent completion.¹³ Hence, this per-stage variability can naturally be exploited to improve average system latency and throughput.

Second, asynchronous pipelines inherently provide *elasticity* (i.e., a variable number of data items can appear in the pipeline at any time, unlike in a synchronous pipeline).^{1,8} If there is no congestion and data items are injected at wide intervals, data items are widely spaced in the pipeline and travel rapidly through. If input rates are higher, spacing becomes tighter between items. In the extreme case, with a slow or stalled output environment, data items become bunched or stalled at close intervals. In all cases, input data items are processed as

they arrive, even with an unknown or irregular arrival rate; there is no wait for a clock edge. Hence, the intertoken spacing and the throughput rate are determined dynamically.

Third, asynchronous pipelines automatically provide flow control. A handshaking protocol inherently offers underflow and overflow protection, even in variable-speed environments. In contrast, synchronous pipelines by default include no flow control. Synchronous flow control is typically supported using explicit credit-based techniques involving extra registers¹⁴ or complex decoupled latch control.¹⁵ A stall signal, used for back pressure, must also be synchronized to the clock at every stage.

Finally, asynchronous pipelines consume dynamic power only on demand. That is, switching activity occurs only when data items are being processed; otherwise, stages and their control are quiescent. Effectively, this approach achieves the benefit of automatic clock gating, but without extra instrumentation, at arbitrary levels of granularity in the design, including under rapidly changing traffic patterns. Furthermore, asynchronous pipelines inherently obviate the need for global clock distribution.

Linear and nonlinear pipeline structures

The focus of this article is on linear pipelines, in which each stage has a single input channel and a single output channel. However, to construct complex digital systems with varied topologies, various other pipeline components are needed, such as parallel forks and joins, conditional splits and merges, arbitration, and loop control structures. These components typically involve small extensions to the linear stages, and their designs are covered extensively in the literature.^{2,3,6-8,16,17}

Benefits of asynchronous pipelines

The preceding four features make asynchronous pipelines highly attractive as a medium for assembling scalable complex systems. An important goal is to design a pipeline with a throughput comparable to a high-performance synchronous pipeline, but with lower end-to-end latency, dynamic power that naturally adapts to the actual traffic, and the flexibility to handle variable input and output rates and variable-delay stages. As a result, an asynchronous pipeline can be reused, without

modification, when interfacing to a large variety of input and output environments. It can also gracefully support dynamic voltage and frequency scaling (DVFS) at its interfaces.

In addition, if mixed-timing interfaces¹⁸ or stoppable clocks⁹ are used, asynchronous pipelines can connect arbitrary, unrelated synchronous clock domains. Hence, these pipelines can serve as the foundation for designing globally asynchronous, locally synchronous (GALS) systems. A GALS approach has been used for several recent pipelined networks on chips (NoCs).^{11,16,19} Several pipelined asynchronous NoCs have been shown to provide significantly lower dynamic power and system latency than comparable single-clock¹⁶ and multiclock²⁰ NoCs.

Finally, there is increasing interest in using asynchronous pipelines as the organizing structure for emerging technologies, such as quantum-dot cellular automata (QCA)²¹ and carbon nanotubes (CNTs), whose timing irregularities make fixed-rate clock distribution extremely difficult.

Asynchronous communication protocols

As Figure 2 shows, there are two basic asynchronous signaling protocols: four-phase (i.e., return-to-zero [RZ])^{19,22,23} and two-phase (i.e., non-return-to-zero [NRZ]).^{1,2,9}

Four-phase handshaking

In four-phase handshaking, *req* (request) and *ack* (acknowledge) are initially deasserted low. The sender initiates a transaction with a rising *req*, and the receiver responds with a rising *ack*. These two events constitute the *evaluate phase*. The two signals, in turn, are then deasserted low in the *reset phase*. Although the protocol requires two round-trip communications per transaction, it is widely used.^{4,6,8-11,19,22,23} The benefit of returning all signals to a baseline state (i.e., all 0) between transactions is simpler hardware design. In addition, the performance overhead of the reset phase can sometimes be reduced or eliminated by overlapping it with other useful operations.

Two-phase handshaking

In two-phase handshaking (i.e., transition signaling) a single request *transition* (or toggle) from the sender initiates a transaction, which is followed by a single acknowledge transition from the receiver.

The protocol requires only one round-trip communication per transaction. Although hardware in some cases is more complex, the throughput and power benefits of a single round-trip communication per transaction are significant, especially for long-distance communication. This approach is also widely used.^{1,2,9,12,16}

Asynchronous data encoding

To encode data, the single control *req* wire is replaced by a data field. Figure 3 shows two common approaches, each for four-phase (RZ) handshaking.

Dual-rail codes and delay-insensitive encoding

In dual-rail encoding,^{8,9,22} each bit is implemented by a pair of wires, or *rails*, as shown in Figure 3a. In the reset phase, both rails are low, forming a *spacer*, which indicates no valid data. During the evaluate phase, the sender asserts exactly one of the two rails high, thereby indicating the actual value of the bit (0 or 1) as well as data validity. The receiver typically uses a completion detector (CD) to identify that a valid codeword has been received. The remainder of the fourphase protocol proceeds as expected: after the dual-rail data is transmitted, *ack* is asserted high by the receiver, the dual-rail *data* is

reset to zero, and *ack* is deasserted low by the receiver. Such an encoding is one instance of a *delay-insensitive (DI) code*,²² in which each codeword uniquely identifies its own validity.

Single-rail bundled data

An alternative encoding scheme is single-rail bundled data (Figure 3b).^{1,9} This scheme uses a

synchronous-style single-rail data channel, called a data bundle. One extra wire is added, a bundling signal (*req*), which uses a worst-case delay and is guaranteed to arrive at the receiver after all data bits are stable and valid. This signal serves as a local strobe, which is passed from sender to receiver. A four-phase protocol is used on the *req* and *ack* signals. Interestingly, data may exhibit glitches arbitrarily

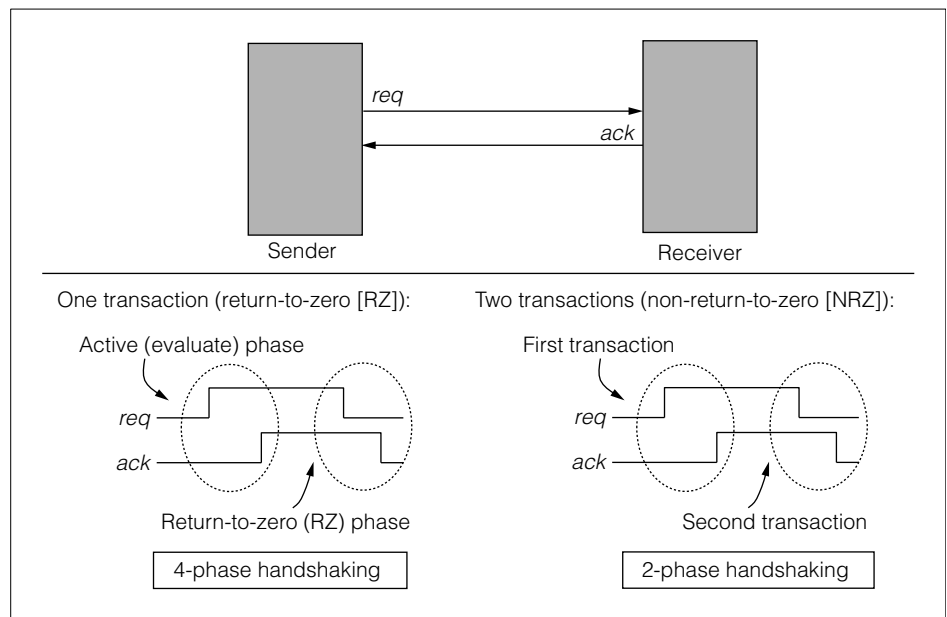


Figure 2. Basic asynchronous communication protocols: four-phase and two-phase.

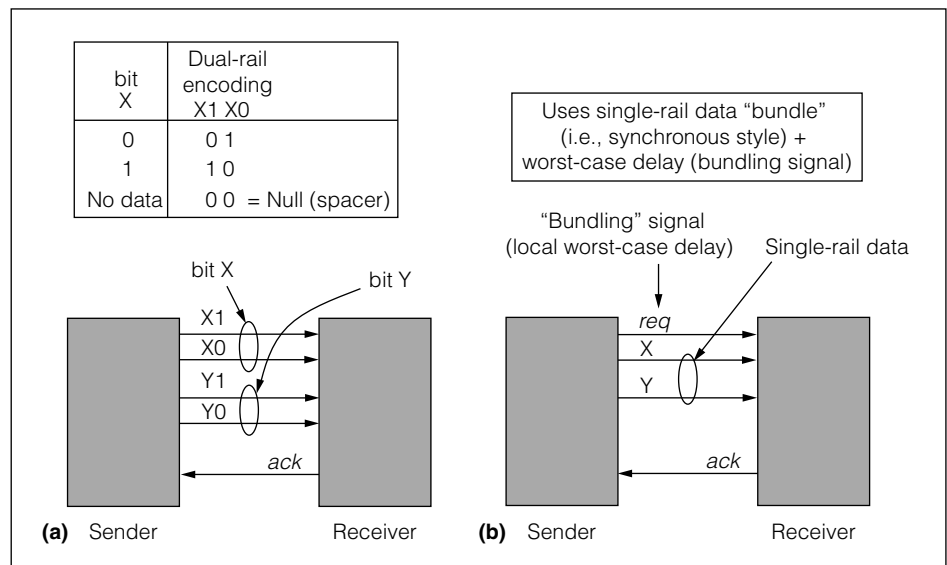


Figure 3. Asynchronous data encoding: dual-rail encoding (a) and single-rail bundled data (b), both four-phase (RZ).

between transactions, as long as it is stable and valid at a setup time before the next *req* is asserted high. Data typically must remain stable from this time until a hold time after *ack* is deasserted low, although some schemes let data change after *ack* is asserted high. These setup and hold time constraints can be satisfied by adding delays, if necessary, on the *req* and *ack* signals. Two-phase versions of single-rail bundled data are also commonly used.^{1,2}

One benefit of a bundled-data scheme is that arbitrary implementations of combinational function blocks can be safely used, including standard synchronous library components, even those that might have hazards. The delay matching of the bundled signal is typically implemented via an inverter chain or a carefully replicated critical path taken from the function block. Although this approach requires worst-case timing, it is a localized constraint. In particular, unlike synchronous design, the stages can be unbalanced, each with its own matched delay. Moreover, the timing margins tend to be fairly tight because some parameters (e.g., process, voltage, and temperature) tend to be locally more uniform.

Trade-offs

Although dual-rail encoding has the overhead of decreased coding efficiency (i.e., double the number of wires), it facilitates timing-robust communication: bits can be skewed arbitrarily because of static (process variation, routing) or dynamic (inductance, crosstalk, thermal) timing deviations, and the receiver can still uniquely identify when a complete code-word is received. These codes are widely used, both for four-phase^{4-6,10,19,22} and two-phase²⁴ protocols. Single-rail bundled data is also widely used.^{1-3,7,9,16,17} Although it has the added design requirement that the bundling signal must always arrive later than the data bundle, it also offers high coding efficiency, low dynamic power, and ease in incorporating existing synchronous function blocks.

Asynchronous pipeline performance

Unlike a synchronous pipeline, in which each stage operates at a fixed clock rate, an asynchronous pipeline operates in a distributed manner, such that each stage's maximum performance is defined only by local parameters.

Two key metrics characterize the performance of an individual asynchronous pipeline stage. The *forward latency* is the time it takes one data item to

flow through a stage, assuming that stage was empty and ready. The *reverse latency* is the time it takes a "hole" (or *ack*) to flow backward through a stage, assuming the stage was initially full. Typically, a complete cycle, defining a stage's maximum throughput, consists of one forward latency (data moving into the next stage) and one reverse latency (*ack* received from the next stage); hence, the cycle time is the sum of the two latencies. More details are available elsewhere in the literature.^{2,4,5,25}

Static logic pipelines

The first class of asynchronous pipelines operates on static logic data paths. Each approach uses single-rail bundled data, as well as explicit storage latches to separate adjacent stages.

Sutherland's micropipeline

Figure 4 shows a basic example of Sutherland's micropipeline, which uses a two-phase protocol.¹ This design and protocol serves as a point of departure for several more-advanced approaches. (Muller presents a four-phase version, which uses a similar control structure.⁸)

Pipeline structure. The pipeline has three components: data (i.e., channels and logic blocks), control, and latches. The leftmost channel has single-rail data ($data_{in}$) and a bundling signal (req_0) as input, and an acknowledgment (ack_1) as output. The rightmost channel has a similar single-rail bundled interface. The *req* and *ack* signals are also used as part of the control component. Each stage must preserve the bundling constraint: a delay element is added to each req_i to match or exceed the worst-case path through the corresponding logic block.

The control provides local clocking. It consists of a simple chain of Muller C-elements, a common asynchronous sequential component.^{1,8} In a two-input C-element, if both inputs are 0, the output is 0, and if both inputs are 1, the output is 1; otherwise, the output maintains its previous value. In this design, the rightmost input of each C-element is inverted.

Storage is designed via specialized *capture-pass latches*, which use transition-based control signals but provide transparent latch operation. Each latch has two control inputs (*C* and *P*, capture and pass) and two control outputs (*C_d* and *P_d*, capture done and pass done). Initially, all signals are 0,

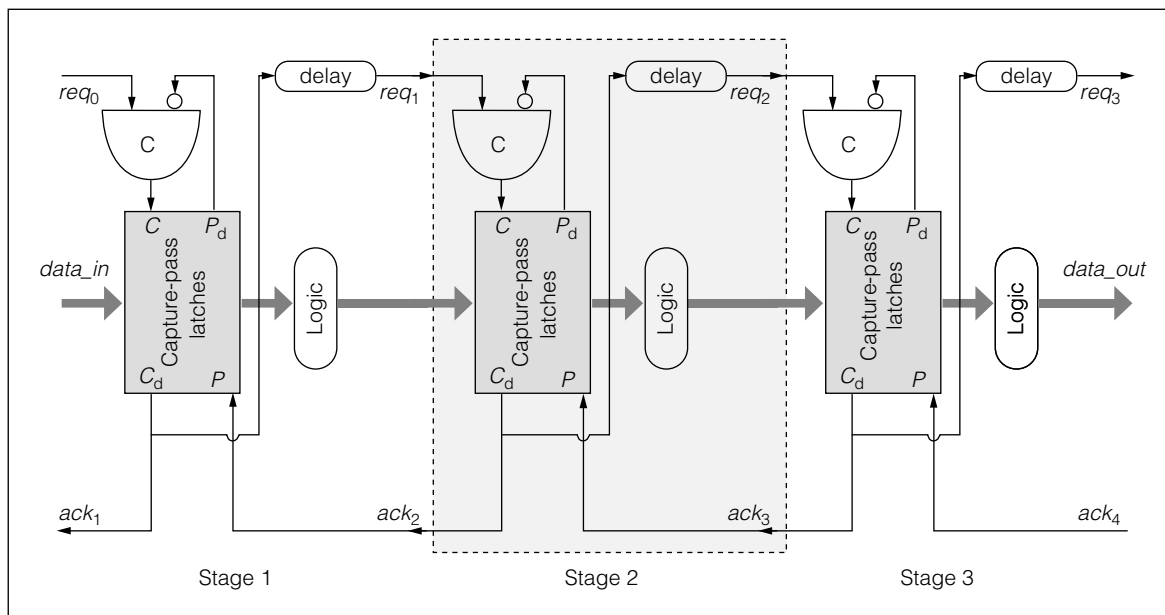


Figure 4. Sutherland's micropipeline using a static logic data path. (C: capture; C_d : capture done; P: pass; P_d : pass done.)

and the latches are transparent. A wire transition from 0 to 1 first arrives at the capture input C , and the latches become opaque; the capture output C_d then makes a transition from 0 to 1 after the hold time is satisfied. Next, when a 0-to-1 transition arrives at the pass input P , the latches become transparent again; the pass output P_d makes a transition from 0 to 1 after the hold time is satisfied. As a result, these signals are all at 1, and the transaction is complete. In this two-phase protocol, for the next transaction, an identical operation occurs, but where the inputs and outputs will make the reverse transitions from 1 to 0.

Pipeline operation. A key feature of this pipeline is that, unlike almost all synchronous latch-based pipelines, the latches are all normally transparent: data can pass directly through them. Hence, initially, the entire pipeline forms a flow-through combinational path. Locally, after data advances through an individual stage's latches, a transition on the C control input makes those latches opaque, thereby storing and protecting the data from any further changes on the stage's input channel. Once data advances through the next stage's latches, a transition on the P control input, communicated from the next stage, makes the current stage's latches transparent, thus safely allowing the next data item to enter. The result is a so-called *capture-pass protocol*.

As an example, consider a simple simulation. Assume an initially empty pipeline in which all latches are transparent and there is no congestion. All C , C_d , P and P_d signals are initially 0. Each C-element in the control chain is therefore half-enabled: with a 0 for the left input and a 1 for the right input after inversion, and with a 0 output. The data first arrives at the left channel and passes directly through all the latches to the rightmost channel. At each stage, in turn, the corresponding req_i bundling signal is toggled slightly after the data has passed through its latches. As a result, the C-element's output toggles to 1, thereby capturing the data (i.e., making the latches opaque). This operation is replicated at each stage; we refer to it as *forward synchronization*.

Concurrently, once data is captured at any stage, the stage signals to its predecessor through a transition on its C_d output (i.e., ack_i) to 1. This event makes the predecessor's latches transparent again (i.e., a P transition), thereby enabling the advance of the next data item. We refer to this operation as *backward synchronization*.

For the next data item, a similar sequence of events occurs, but with all signals toggling in the opposite direction. Effectively, each data item initiates a "wave front," which advances through the pipeline and is protected by a series of latch-capture operations. Predecessor stages, behind the wave front, are subsequently freed up through a series of pass

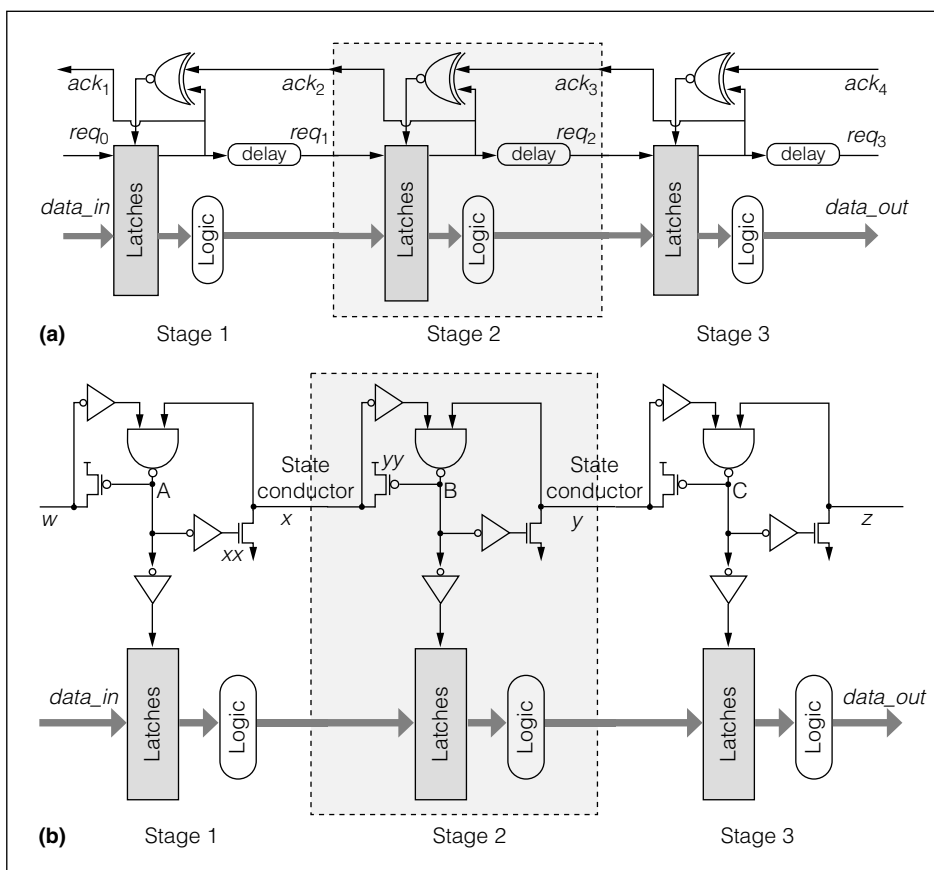


Figure 5. Recent high-performance pipeline styles with static logic data paths: Mousetrap pipeline (a) and GasP pipeline (b).

operations, once data has been safely copied to the next stage. The old data can then be overwritten by the next wave front.

This protocol allows dynamically variable spacing of data items in the pipeline, depending on their input arrival rate and the logic block delays. Congestion is communicated to the left by withholding of ack_i transitions. In this case, the backward chain of pass operations is interrupted, and each stage stalls, storing a distinct data item.

Mousetrap pipeline

We developed Mousetrap at Columbia University to be a high-performance pipeline that supports the use of a standard cell methodology.² Although its capture-pass protocol is based on that of micropipelines, it has less complex signaling and far lower overhead, and it uses simpler control and data latches.

Figure 5a shows a basic Mousetrap pipeline. Like a micropipeline, it uses a two-phase protocol. The local control for each stage is only a single combinational

exclusive-NOR (XNOR2) gate. The capture-pass latches of micropipelines are replaced by a single bank of level-sensitive D-latches, which are available in standard cell libraries. The result is a lightweight control and storage structure.

The forward synchronization operation is similar to that of micropipelines, but with a different realization. Initially, suppose all req_i and ack_i signals are at 0, and local latch controls (i.e., XNOR2 outputs) are at 1. All data latches are therefore transparent, forming a flow-through combinational system. Once data arrives at the left channel, it passes directly through the entire series of data latches to the right channel. At each stage, after the data arrives and passes through its latches, the corresponding req_i bundling signal toggles from 0 to 1. As a result, the stage's XNOR2 control output toggles from 1 to 0 and captures the data, thereby protecting

it from being overwritten by any new data item from the left neighbor.

At the same time, when data is about to be captured at any stage i (i.e., the req_i output of its latches has toggled from 0 to 1 and is enabling the deassertion of the XNOR2 control output), the stage requests the next data item from its left neighbor. In particular, stage i signals to its predecessor, stage $i - 1$, through a transition on its ack_i output from 0 to 1. This enables the pass operation of stage $i - 1$, thereby causing the latches of stage $i - 1$ to become transparent. This backward synchronization effectively indicates that stage i is in the process of safely storing the current data item, and informs stage $i - 1$ that its output will no longer be needed and can be overwritten. Likewise, stage i itself enters its pass phase when stage $i + 1$ indicates its data is being stored. Subsequently, when the next data item enters the pipeline, the same two-phase protocol is repeated, but with the bundling signal at each stage toggling in reverse from 1 to 0.

For correct operation, a simple one-sided hold-time constraint is required: a stage must fully capture its current data (i.e., all data latches become opaque) before its predecessor has passed any new data to its latch inputs. This constraint can always be satisfied by adding sufficient delay to the backward ack_i output, if necessary.²

Interestingly, unlike almost all synchronous designs using single data latches, Mousetrap has no two-sided timing constraints (e.g., no pulse mode, or short versus long path constraints). In addition, in congested scenarios, every stage can hold a distinct data item, thereby providing 100% storage capacity with only single data latches separating adjacent stages.

GasP pipeline

Finally, another approach, GasP, was developed at Sun Research Laboratories.³ Like Mousetrap, GasP uses single D-latches for the data path. However, the pipeline control includes dynamic logic and custom gates, along with careful path timing. Figure 5b shows a basic GasP pipeline.

GasP has several interesting features. First, unlike Mousetrap, the data latches are normally opaque, and a pulse-based latch control protocol is used. Second, each control channel between adjacent stages consists of a single wire (indicated as a state conductor in the figure) instead of the usual pair of request and acknowledge wires. This single-track wire channel passes communication in both directions. Initially, the wire is deasserted high. To initiate a communication, an active-low request is sent from left to right. To complete the communication, an active-high acknowledgment is sent from right to left. As a result, the channel returns to its default high state. No more than one channel driver, left or right, can be active at any time, and the control circuit is locally timed to avoid transient short circuits. This novel protocol eliminates the need for two wires; it effectively multiplexes the request and acknowledge onto a single wire.

Hence, GasP effectively combines the benefits of both two-phase and four-phase protocols. Like a two-phase protocol, there is only one round-trip communication through GasP's single-track channel. However, as in a four-phase protocol, the channel always returns to the same value whenever a transaction is complete.

Despite the protocol differences, the forward synchronization operation of this pipeline is

fundamentally the same as in the previous designs. If the pipeline is empty, all single-track channels (w, x, y, z) are initially deasserted high, and data-latch enable signals are deasserted low. When a left request arrives by asserting w low, this request causes each successive single-track channel to be driven active-low through the inverters and the NAND gate in each GasP stage. The active-low transition on each NAND gate also causes the corresponding data latch enable signal to be asserted high, thereby driving the stage's data latches to become transparent. As an immediate result, signal transitions on the two short local loops deassert both inputs of each NAND2 gate low, thereby releasing the single-track channels and deasserting low the latch enable signals. The net result is that each stage's latch enable undergoes a short pulse, allowing the data to propagate forward through the pipeline stage. In the reverse synchronization operation, backward acknowledgments are driven high on each channel, which is thus reset to its initial value.

The circuit designs are highly optimized for low latency (i.e. use low "logical effort"), require balanced control path delays, and must satisfy two-sided timing constraints (i.e., short and long path requirements). A more aggressive implementation has also been proposed.³

Several variants with longer forward and reverse latencies have also been introduced. In particular, to accommodate longer-latency logic blocks, a designer can insert pairs of inverters on the forward path to delay the firing of the n-transistor. Similarly, a designer can add pairs of inverters on the reverse path to delay the firing of the p-transistor, thereby slowing down the pipeline operation for improved timing margins.

Dynamic logic pipelines

The second class of asynchronous pipelines operate on dynamic logic data paths. All the approaches reviewed here use a four-phase (RZ) protocol and are latchless (i.e., no explicit storage elements are needed between adjacent stages).

Trade-offs in using dynamic logic

Dynamic data paths are common in high-performance digital systems. By eliminating complex pull-up transistor networks, dynamic gates can provide the benefits of reduced chip area and

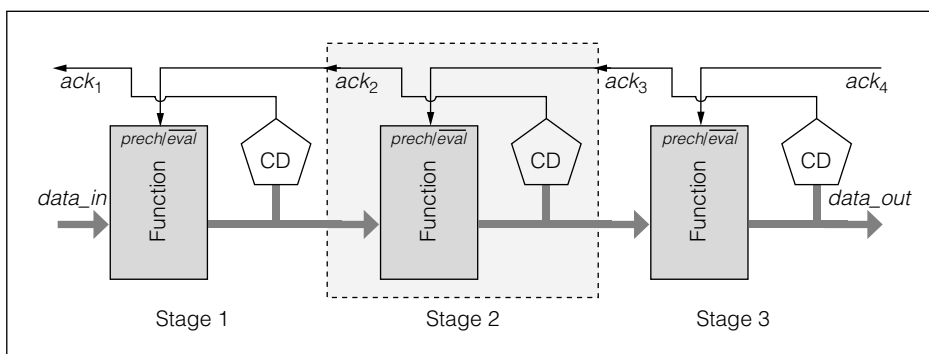


Figure 6. Classic dynamic-pipeline style: Williams and Horowitz' PS0 pipeline. (CD: completion detector.)

reduced switched capacitance, which in turn can lead to higher speed and lower energy consumption. However, dynamic logic also has its drawbacks: greater design and validation effort, and less noise immunity. Therefore, it is typically used only in speed-critical parts of a design—for example, in ASICs and in arithmetic and logic units used in high-speed microprocessors. However, recent industrial efforts continue to demonstrate the viability and benefits of dynamic logic, even in modern advanced VLSI processes. For instance, Intrinsity (recently acquired by Apple) has developed a complete general-purpose synchronous CAD design flow based on domino logic, and used it to implement high-performance and low-power processor cores, including an ARM Cortex-A8 processor in 45-nm technology.

For several reasons, dynamic logic is an especially good match for asynchronous pipelines. In particular, in dynamic asynchronous systems, the gate type is typically *fully staticized domino*,^{4-7,11,17} in which each dynamic gate output has an attached primary inverter along with a weak feedback inverter, forming a lightweight storage element. Such gates have good immunity to the effects of leakage, charge sharing, and noise. Moreover, DI-encoded asynchronous data paths can gracefully accommodate delay variations introduced by uncertainty in charge sharing and noise, because individual bits can safely arrive with arbitrary skew. As a result, dynamic logic has been widely used in several recent high-performance asynchronous commercial products, such as Ethernet switch chips from Fulcrum Microsystems (currently 65 nm) and FPGAs from Achronix Semiconductor (90-nm down to 22-nm technologies).

Dynamic logic and asynchronous pipelines

A unique feature of many asynchronous dynamic pipelines is that they are latchless. Thus, with proper sequencing of control operations, an asynchronous pipeline can exploit the implicit latching functionality of dynamic gates and entirely avoid explicit storage elements between stages. Achieving similar latchless operation in a synchronous implementation

typically would require using complex multiphase clocking. This asynchronous feature provides the benefits of reduced critical delays, smaller chip area, and lower power consumption, thereby minimizing some of the key overheads of fine-grained pipelining.

Several asynchronous dynamic logic approaches have been proposed.^{4,6,7,17,23,26,27} We will begin by reviewing the PS0 pipeline style by Williams and Horowitz,⁴ which is influential and an important foundation for most later styles. The “Advanced Asynchronous Dynamic Pipeline Styles” sidebar highlights two recent higher-performance approaches: a timing-robust style called the precharge half-buffer (PCHB) by Lines,⁶ and our high-capacity (HC) style,⁷ which provides high throughput and high storage density.

The PS0 pipeline style

The PS0 pipeline style was introduced to support pipelined computation in which low latency is critical.⁴ Because the pipeline has no latches, there is no sequential overhead (i.e., no latches or registers) on its forward path. As a result, in iterative computations that are implemented using self-timed rings, the zero overhead to latency results in purely combinational execution times for multistage operations. This benefit was exploited in the design of self-timed floating-point dividers based on the PS0 pipeline style in two commercial processors developed at HAL Computers in the 1990s.⁵

Figure 6 shows the basic structure of a PS0 pipeline; each stage consists of a function block and a CD. The data path uses DI coding (in particular, dual-rail encoding) and contains only function blocks, without any explicit registers.

Each function block alternates between an evaluate phase and a precharge phase. Initially, the function blocks have all 0 outputs, and hence are reset. In the evaluate phase, the function is computed after its data inputs arrive. In the precharge phase, the function block is reset, with all its outputs returning to 0. An important property of dynamic logic in the precharge phase is that each stage has an inherent blocking capability: new inputs applied to its function are ignored by the dynamic block, and the stage's outputs remain reset. As a result, evaluation and precharge of a dynamic function block are somewhat analogous to making a latch transparent and opaque, respectively, in static logic pipelines.

A CD is attached to each stage's output channel. In each evaluate phase, the CD indicates when a complete DI codeword has been generated (i.e., a function block's computation is complete). In each precharge phase, the CD indicates when a function block's outputs have all been reset to 0. The single precharge-and-evaluate control input, *prech/eval*, of a stage's function block is connected from the output of the next stage's CD. A high value causes precharge, and a low value enables evaluation. No latches are needed between function blocks.

The coordination of pipeline control in a PS0 pipeline is quite simple, with two backward synchronization events per cycle. These events follow two basic rules: a stage is precharged whenever the next stage finishes evaluation, and a stage evaluates whenever the next stage finishes its precharge. This protocol ensures that each pair of consecutive data tokens is always separated by a reset token or spacer, in which each dual-rail data bit in a stage is reset to a 00 value.

As an example, consider a simple simulation of the PS0 pipeline operation. We can derive the complete cycle of events for a PS0 pipeline stage by observing how a single data item propagates after arriving at the left input channel of an initially empty three-stage pipeline. Initially, all pipeline stages are in the evaluate phase, ready to receive new data. After the data input arrives at stage 1, there are three events from the start of the stage's evaluation phase to the start of its subsequent precharge phase. First, stage 1 evaluates, and produces a valid output. Second, stage 2 evaluates, and produces a valid output. Third, stage 2's completion detector detects completion of its evaluation.

Finally, stage 1 is enabled by the pipeline control to enter its precharge phase.

Extending this simulation, we can trace the entire cycle time of stage 1, from the start of its current evaluation phase, through its precharge phase, to the start of its next evaluation phase. There are six successive events:

1. Stage 1 evaluates.
2. Stage 2 evaluates.
3. Stage 3 evaluates.
4. The CD of stage 3 detects completion of evaluation and initiates the precharge of stage 2.
5. Stage 2 precharges.
6. The CD of stage 2 detects completion of precharge, thereby releasing the precharge of stage 1 and enabling stage 1 to evaluate once again.

Thus, the cycle time of a pipeline stage is the sum of the delays associated with these six events.

For a dual-rail code, a CD is implemented as follows (details are available elsewhere^{4,5,9}): Each pair of rails, forming a bit, is the input to a distinct OR2 gate, which evaluates if either rail has been asserted high. The outputs of all such OR2 gates are then combined together through a tree of C-elements.^{1,4,8} Each C-element's output is asserted high when all its inputs are high, and asserted low when all its inputs are low; otherwise, it holds its current output value. Hence, the CD asserts its output high when all bits have arrived, and asserts its output low when all bits have been reset.

The distributed nature of asynchronous control is key to obtaining latchless operation in a PS0 pipeline. In particular, the pipeline style ensures that, once stage 1 has completed evaluation, it will hold its output stable as long as stage 2 is still using the data of stage 1 (i.e., until stage 2 sends stage 1 an acknowledgment). Thus, with this local control sequencing, the dynamic function block itself implicitly provides the storage functionality needed. In contrast, typical synchronous pipelines employ a two-phase clock, which causes stage 1 to start resetting while stage 2 starts evaluating. This race condition in a synchronous pipeline requires the insertion of a latch between the two function blocks; alternatively, complex multiphase synchronous clocking would be necessary to achieve similar effects.

Advanced Asynchronous Dynamic Pipeline Styles

Several alternative dynamic pipeline styles have been proposed specifically for achieving higher throughput.¹⁻⁶

Andrew Lines introduced the precharge half-buffer (PCHB) style for the robust design of asynchronous pipelines in the presence of timing variations (see Figure A).¹ This style uses *quasi-delay-insensitive* (QDI) control circuits, meaning its functionality is correct even in the presence of arbitrary gate and wire delays, provided only that delays along all the branches of each wire fork are nearly equal (the so-called *isochronic fork* assumption). A key advantage of QDI styles is that they simplify the task of timing verification, thereby greatly alleviating designer effort. It also uses DI encoding for its data path.

A PCHB stage has two completion detectors: one on the input side and one on the output side, unlike PS0 pipelines. As a result, also unlike a PS0 stage, a PCHB stage does not start evaluation until all input bits are valid, as indicated by its input completion detector (CD). It thereby absorbs, rather than propagates, skew across individual bits in the data path. After the inputs are received, the complete cycle of events for a PCHB stage is otherwise quite similar to that of PS0. These pipelines have been successfully used in commercial high-speed communications chips designed by Fulcrum Microsystems.⁷

We introduced another dynamic-logic style, the high-capacity (HC) pipeline,⁴ to double the storage density and increase throughput over that of PS0

(see Figure B). The pipeline uses single-rail bundled data.

A drawback of most latchless asynchronous pipelines, including PS0 and PCHB, is that they provide only 50% storage capacity, in which consecutive data items must be separated by a stage with a spacer. As a result, in applications where buffering capacity is important, deep pipelines must be implemented. In contrast, under congestion, each HC stage can hold a distinct data item, thereby providing 100% storage capacity.

HC pipelines also use a fundamentally different structure and protocol than PS0. In particular, its dynamic gates are controlled by two decoupled inputs, *prech* and *eval*. These inputs provide a novel *three-phase operation*: precharge, evaluate, and isolate. Precharge occurs when *prech* is asserted low, and evaluation occurs when *eval* is asserted high. When both signals are deasserted, the gate enters an isolate phase, in which its output is effectively in a hold state, immune from all gate inputs; pull-up and pull-down are then simultaneously inactive and the gate can neither evaluate nor precharge. The isolate phase is the key to this protocol: because a stage quickly isolates itself when finishing evaluation, the previous stage can immediately proceed with not only its precharge but also the subsequent evaluation of the next data item. The result is significantly greater concurrency, yielding both higher

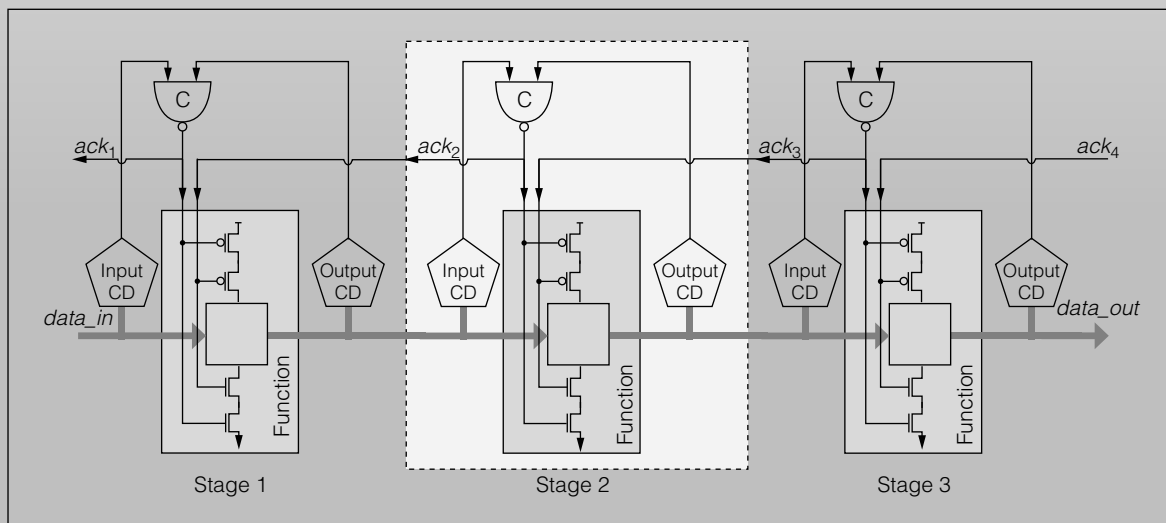


Figure A. Lines' precharge half-buffer (PCHB) pipeline. (CD: completion detector.)

Two other dynamic styles have also been proposed that feature an isolate phase and provide 100% storage capacity, including a domino version of GasP⁵ (based on HC), and a relative of PCHB called precharge full buffer or PCFB.¹ However, unlike HC pipelines, domino GasP includes two-sided timing requirements and the use of single-track channels. The PCFB style offers lower throughput than HC, and is instead targeted for timing-robust QDI operation. Interestingly, while PCHB also offers an isolate phase, its protocol is less concurrent and therefore provides only 50% storage capacity.

References

1. A.M. Lines, *Pipelined Asynchronous Circuits*, tech. report, Dept. of Computer Science, California Inst. of Technology, 1998.

2. S.B. Furber and J. Liu, "Dynamic Logic in Four-Phase Micro-pipelines," *Proc. 2nd Int'l Symp. Advanced Research in Asynchronous Circuits and Systems (ASYNC 96)*, IEEE CS Press, 1996, pp. 11-16.
3. M. Singh and S.M. Nowick, "The Design of High-Performance Dynamic Asynchronous Pipelines: Lookahead Style," *IEEE Trans. Very Large Scale Integration (VLSI) Systems*, vol. 15, no. 11, 2007, pp. 1256-1269.
4. M. Singh and S.M. Nowick, "The Design of High-Performance Dynamic Asynchronous Pipelines: High-Capacity Style," *IEEE Trans. Very Large Scale Integration (VLSI) Systems*, vol. 15, no. 11, 2007, pp. 1270-1283.
5. J. Ebergen et al., "GasP Control for Domino Circuits," *Proc. 11th IEEE Int'l Symp. Asynchronous Circuits and Systems (ASYNC 05)*, IEEE CS Press, 2005, pp. 12-22.
6. M. Ferretti and P.A. Beerel, "High Performance Asynchronous Design Using Single-Track Full-Buffer Standard Cells," *IEEE J. Solid-State Circuits*, vol. 41, no. 6, 2006, pp. 1444-1454.
7. A. Lines, "Asynchronous Interconnect for Synchronous SoC Design," *IEEE Micro*, vol. 24, no. 1, 2004, pp. 32-41.
8. M. Singh et al., "An Adaptively Pipelined Mixed Synchronous-Asynchronous Digital FIR Filter Chip Operating at 1.3 Gigahertz," *IEEE Trans. Very Large Scale Integration (VLSI) Systems*, vol. 18, no. 7, 2010, pp. 1043-1056.



More recently, an alternative family of look-ahead pipelines (LPs) provide optimizations that accelerate the PS0 computation, and therefore increase throughput, for both single-rail bundled and dual-rail data paths.¹⁷

Choosing an asynchronous pipeline style

With many asynchronous pipeline styles available, a designer must consider several criteria to choose an approach that is a good match for the intended application. The choice of static versus dynamic data paths is typically determined by the design tools and design effort available. In particular, static data paths are more amenable to synthesis using standardized design tools and standard cell libraries. Dynamic logic, on the other hand, requires custom gates, as well as specialized design and analysis tools for verifying timing and noise immunity. On the whole, dynamic logic requires significantly greater designer effort, although it can yield somewhat higher performance. Similarly, single-rail bundled-data design generally requires greater design and timing validation effort than DI encoding to ensure bundling constraints are met, but it has the benefits of smaller area, lower power, and ease of reuse of synchronous function blocks. DI encoding techniques can lead to 2 to 4 times as much switching activity in the data path²⁸—although, recently, *m-of-n*²² and level-encoded transition signaling (LETS)²⁹ codes have been shown to greatly reduce this gap. However, if timing-robust operation in the presence of high variability is paramount, then using DI coding becomes imperative, as the timing margins needed for safe delay matching with single-rail bundled data become unwieldy.

THE SIX PIPELINE STYLES reviewed in this article represent different trade-offs between performance, power, and ease of design. Of the three static pipelines, the GasP approach offers the highest performance, although it involves significant design effort because of its complex circuit structure and operation, as well as its stringent timing constraints, and is therefore more difficult to use with automated synthesis flows. The Mousetrap approach is next in performance, and has the added benefit of an entirely standard cell implementation; it is therefore well-suited for automation. Of the

three dynamic pipelines, PS0 is the simplest to implement, but it offers the lowest throughput. The HC style has the highest throughput and lowest energy consumption of the three, but it requires matching of bundling delays. PCHB uses quasi-delay-insensitive (QDI) control logic and DI coding of the data path, and hence is highly robust to variability.

Recent commercial and experimental chips have validated the high performance of the asynchronous pipeline styles reviewed in this article. In particular, a test chip for a fine-grained pipelined greatest common divisor (GCD) computation has demonstrated Mousetrap pipelines capable of operating at 2.1 GHz in 130-nm technology.³⁰ GasP pipelines have demonstrated very high performance, owing to their low-overhead protocol,³ including recent results of 4 GHz in 90-nm technology for the Infinity test chip at Sun Labs (now Oracle Labs). Fulcrum Microsystems' commercial Nexus crossbar switch, operating at 1.35 GHz in 130-nm technology,¹¹ uses the PCHB pipeline style to implement the data path.⁶ Finally, in an experimental, mixed-synchronous/asynchronous FIR filter chip developed at IBM Research, operating at 1.8 GHz in 180-nm technology,¹² the HC style is used to implement the speed-critical data path.⁷

Several support tools, design flows, extensions, and applications for high-performance asynchronous pipelines have been explored. These include system-level performance analysis²⁵ and optimization^{31,32} techniques, as well as some initial automated CAD synthesis flows.^{24,27} Low-overhead testing techniques have also been proposed.^{33,34} Recent pipelined applications include high-performance FPGAs,¹⁰ Ethernet switch chips,¹¹ iterative dividers,^{4,5} FIR filters,¹² and NoCs.^{16,19}

Synchronous methodologies that borrow from the asynchronous design approach have also been proposed,^{14,15} which adopt asynchronous ideas of back pressure and handling of long global paths within clocked systems. In addition, a desynchronization methodology synthesizes asynchronous circuits from a synchronous netlist by replacing the clock with handshaking channels.³⁵

Thus, there are exciting advances in the design, optimization, tool support, and application of high-speed asynchronous pipelines. There is also increasing cross-fertilization with, and migration of asynchrony into, the synchronous world. ■

Acknowledgments

We appreciate the funding support of the National Science Foundation under grants CCF-0964606, CCF-0811504, and CCF-0702712. We thank Jordi Cortadella (Polytechnic University of Catalonia, Spain), Gennette Gill (Columbia University), Andrew Lines (Fulcrum Microsystems), Rajit Manohar (Cornell University), Jens Sparsoe (Technical University of Denmark), and Ivan Sutherland (Portland State University) for their technical suggestions.

References

1. I.E. Sutherland, "Micropipelines," *Comm. ACM*, vol. 32, no. 6, 1989, pp. 720-738.
2. M. Singh and S.M. Nowick, "MOUSETRAP: High-Speed Transition-Signaling Asynchronous Pipelines," *IEEE Trans. Very Large Scale Integration (VLSI) Systems*, vol. 15, no. 6, 2007, pp. 684-698.
3. I. Sutherland and S. Fairbanks, "GasP: A Minimal FIFO Control," *Proc. 7th Int'l Symp. Asynchronous Circuits and Systems (ASYNC 01)*, IEEE CS Press, 2001, pp. 46-53.
4. T.E. Williams, "Self-Timed Rings and Their Application to Division," doctoral dissertation, Dept. of Electrical Eng., Stanford Univ., 1991.
5. T.E. Williams and M.A. Horowitz, "A Zero-Overhead Self-Timed 160ns 54b CMOS Divider," *IEEE J. Solid-State Circuits*, vol. 26, no. 11, 1991, pp. 1651-1661.
6. A.M. Lines, *Pipelined Asynchronous Circuits*, tech. report no. CaltechCSTR:1998.cs-tr-95-21, Dept. of Computer Science, California Inst. of Technology, 1998.
7. M. Singh and S.M. Nowick, "The Design of High-Performance Dynamic Asynchronous Pipelines: High-Capacity Style," *IEEE Trans. Very Large Scale Integration (VLSI) Systems*, vol. 15, no. 11, 2007, pp. 1270-1283.
8. D.E. Muller, "Asynchronous Logics and Application to Information Processing," *Proc. Symp. the Application of Switching Theory to Space Technology*, Stanford University Press, 1963, pp. 289-297.
9. C.L. Seitz, "System Timing," *Introduction to VLSI Systems*, C.A. Mead and L.A. Conway, eds., Addison-Wesley, 1980, pp. 218-262.
10. C. La Frieda, B. Hill, and R. Manohar, "An Asynchronous FPGA with Two-Phase Enable-Scaled Routing," *Proc. IEEE Symp. Asynchronous Circuits and Systems (ASYNC 10)*, IEEE CS Press, 2010, pp. 141-150.
11. A. Lines, "Asynchronous Interconnect for Synchronous SoC Design," *IEEE Micro*, vol. 24, no. 1, 2004, pp. 32-41.
12. M. Singh et al., "An Adaptively Pipelined Mixed Synchronous-Asynchronous Digital FIR Filter Chip Operating at 1.3 Gigahertz," *IEEE Trans. Very Large Scale Integration (VLSI) Systems*, vol. 18, no. 7, 2010, pp. 1043-1056.
13. S.M. Nowick et al., "Speculative Completion for the Design of High-Performance Asynchronous Dynamic Adders," *Proc. 3rd Int'l Symp. Advanced Research in Asynchronous Circuits and Systems (ASYNC 97)*, IEEE CS Press, 1997, pp. 210-223.
14. L.P. Carloni, K.L. McMillan, and A.L. Sangiovanni-Vincentelli, "Theory of Latency-Insensitive Design," *IEEE Trans. Computer-Aided Design of Integrated Circuits and Systems*, vol. 20, no. 9, 2001, pp. 1059-1076.
15. J. Carmona et al., "Elastic Circuits," *IEEE Trans. Computer-Aided Design of Integrated Circuits and Systems*, vol. 28, no. 10, 2009, pp. 1437-1455.
16. M.N. Horak et al., "A Low-Overhead Asynchronous Interconnection Network for GALS Chip Multiprocessors," *IEEE Trans. Computer-Aided Design of Integrated Circuits and Systems*, vol. 30, no. 4, 2011, pp. 494-507.
17. M. Singh and S.M. Nowick, "The Design of High-Performance Dynamic Asynchronous Pipelines: Look-ahead Style," *IEEE Trans. Very Large Scale Integration (VLSI) Systems*, vol. 15, no. 11, 2007, pp. 1256-1269.
18. T. Chelcea and S.M. Nowick, "Robust Interfaces for Mixed-Timing Systems," *IEEE Trans. Very Large Scale Integration (VLSI) Systems*, vol. 12, no. 8, 2004, pp. 857-873.
19. E. Beigne et al., "An Asynchronous NOC Architecture Providing Low Latency Service and Its Multi-level Design Framework," *Proc. 11th IEEE Int'l Symp. Asynchronous Circuits and Systems (ASYNC 05)*, IEEE CS Press, 2005, pp. 54-63.
20. A. Sheibanyrad, A. Greiner, and I. Miro-Panades, "Multi-synchronous and Fully Asynchronous NoCs for GALS Architectures," *IEEE Design & Test*, vol. 25, no. 6, 2008, pp. 572-580.
21. M. Choi et al., "Efficient and Robust Delay-Insensitive QCA (Quantum-Dot Cellular Automata) Design," *Proc. 21st IEEE Int'l Symp. Defect and Fault-Tolerance in VLSI Systems (DFT 06)*, IEEE CS Press, 2006, pp. 80-88.
22. W.J. Bainbridge et al., "Delay-Insensitive, Point-to-Point Interconnect Using m-of-n Codes," *Proc. 9th Int'l Symp. Asynchronous Circuits and Systems (ASYNC 03)*, IEEE CS Press, 2003, pp. 132-140.
23. S.B. Furber and J. Liu, "Dynamic Logic in Four-Phase Micropipelines," *Proc. 2nd Int'l Symp. Advanced*

- Research in Asynchronous Circuits and Systems* (ASYNC 96), IEEE CS Press, 1996, pp. 11-16.
24. B. Quinton, M. Greenstreet, and S. Wilton, "Practical Asynchronous Interconnect Network Design," *IEEE Trans. Very Large Scale Integration (VLSI) Systems*, vol. 16, no. 5, 2008, pp. 579-588.
 25. P.A. Beerel and A. Xie, "Performance Analysis of Asynchronous Circuits Using Markov Chains," *Proc. Concurrency and Hardware Design*, LNCS 2549, Springer, 2002, pp. 313-344.
 26. J. Ebergen et al., "GasP Control for Domino Circuits," *Proc. 11th IEEE Int'l Symp. Asynchronous Circuits and Systems* (ASYNC 05), IEEE CS Press, 2005, pp. 12-22.
 27. M. Ferretti and P.A. Beerel, "High Performance Asynchronous Design Using Single-Track Full-Buffer Standard Cells," *IEEE J. Solid-State Circuits*, vol. 41, no. 6, 2006, pp. 1444-1454.
 28. P.A. Beerel and M.E. Roncken, "Low Power and Energy Efficient Asynchronous Design," *J. Low Power Electronics*, vol. 3, no. 3, 2007, pp. 234-253.
 29. P.B. McGee et al., "A Level-Encoded Transition Signaling Protocol for High-Throughput Asynchronous Global Communication," *Proc. 14th IEEE Int'l Symp. Asynchronous Circuits and Systems* (ASYNC 08), IEEE CS Press, 2008, pp. 116-127.
 30. G. Gill et al., "A High-Speed GCD Chip: A Case Study in Asynchronous Design," *Proc. IEEE Computer Society Ann. Symp. VLSI*, IEEE CS Press, 2009, pp. 205-210.
 31. G. Gill and M. Singh, "Automated Microarchitectural Exploration for Achieving Throughput Targets in Pipelined Asynchronous Systems," *Proc. IEEE Symp. Asynchronous Circuits and Systems* (ASYNC 10), IEEE CS Press, 2010, pp. 117-127.
 32. P. Prakash and A.J. Martin, "Slack Matching Quasi Delay-Insensitive Circuits," *Proc. 12th IEEE Int'l Symp. Asynchronous Circuits and Systems* (ASYNC 06), IEEE CS Press, 2006, pp. 195-204.
 33. G. Gill et al., "Low-Overhead Testing of Delay Faults in High-Speed Asynchronous Pipelines," *Proc. 12th IEEE Int'l Symp. Asynchronous Circuits and Systems* (ASYNC 06), IEEE CS Press, 2006, pp. 46-56.
 34. O.A. Petlin and S.B. Furber, "Scan Testing of Micropipelines," *Proc. 13th IEEE VLSI Test Symp.* (VTS 95), IEEE CS Press, 1995, pp. 296-301.
 35. J. Cortadella et al., "Desynchronization: Synthesis of Asynchronous Circuits from Synchronous Specifications," *IEEE Trans. Computer-Aided Design of Integrated Circuits and Systems*, vol. 25, no. 10, 2006, pp. 1904-1921.

Steven M. Nowick is a professor of computer science and chair of the Computer Engineering Program at Columbia University. His research interests include the design and optimization of asynchronous and mixed-timing circuits and systems; networks on chips; low-power and high-performance digital design; CAD tools for asynchronous systems; logic synthesis; and encoding techniques for low-power, delay-insensitive communication. He has a PhD in computer science from Stanford University. He is a Fellow of IEEE.

Montek Singh is an associate professor of computer science at the University of North Carolina at Chapel Hill. His research interests include asynchronous and mixed-timing circuits and systems; CAD tools for design, analysis, and optimization; high-speed and low-power VLSI design; and applications to emerging computing technologies and energy-efficient graphics hardware. He has a PhD in computer science from Columbia University.

■ Direct questions and comments about this article to Steven M. Nowick, Dept. of Computer Science, Room 508, Computer Science Building, Columbia University, New York, NY 10027; nowick@cs.columbia.edu; or Montek Singh, Dept. of Computer Science, Sitterson Hall Room FB234, University of North Carolina, Chapel Hill, NC 27599; montek@cs.unc.edu.



Selected CS articles and columns are also available for free at <http://ComputingNow.computer.org>.